

随机数据压测

```
foretoken bench \  
  --url http://127.0.0.1:8008/v1/chat/completions \  
  --model Qwen3.6-27B \  
  --dataset random \  
  --tokenizer-path /exportr/zxcpu2/shiweijie/cache/models--Qwen--Qwen3.6-  
27B/snapshots/6a9e13bd6fc8f0983b9b99948120bc37f49c13e9 \  
  --min-prompt-length 128 --max-prompt-length 512 \  
  --parallel 4 --number 20 --max-tokens 64 \  
  --rate 5 \  
  --wandb
```


```

● (eval) wshiah@zxcpu4:~/code/zhuting/bench/foretoken/benchmarks$ foretoken bench \
  --url http://127.0.0.1:8008/v1/chat/completions \
  --model Qwen3.6-27B \
  --dataset random \
  --tokenizer-path /export/zxcpu2/shiweijie/cache/models--Qwen--Qwen3.6-27B/snapshots/6a9e1
3bd6fc8f0983b9b99948120bc37f49c13e9 \
  --min-prompt-length 128 --max-prompt-length 512 \
  --parallel 4 --number 20 --max-tokens 64 \
  --rate 5 \
  --wandb
2026-08-10 16:03:02,349 INFO [benchmarks.main]
=====
Foretoken Benchmark
=====
Configuration:
  URL       : http://127.0.0.1:8008/v1/chat/completions
  Model     : Qwen3.6-27B
  Parallel  : [4]
  Number    : [20]
  Rate      : 5 req/s (closed-loop, Poisson pacing)
  Open Loop : False
  Stream    : True
  Dataset   : random (prefix=0, min=128, max=512)

/export/zxcpu1/shiweijie/miniconda3/envs/eval/lib/python3.12/site-packages/transformers/util
s/hub.py:110: FutureWarning: Using `TRANSFORMERS_CACHE` is deprecated and will be removed in
v5 of Transformers. Use `HF_HOME` instead.
  warnings.warn(
2026-08-10 16:03:09 - evalscope - INFO: Using 248060 allowed tokens out of 248077 total toke
ns (excluded 17 special/byte-fallback tokens)
2026-08-10 16:03:09 - evalscope - INFO: Sampling input lengths from [118, 503)
2026-08-10 16:03:16,803 - benchmarks.report.wandb_logger - INFO: W&B logging enabled: projec
t=foretoken-bench name=Qwen3.6-27B_20260810_160309 group=-
Benchmarking: 100%|████████████████████████████████████████| 20/20 [00:09<00:00, 2.18it/s]
2026-08-10 16:03:26,007 - benchmarks.runner.base - INFO: Benchmark finished!
2026-08-10 16:03:26,010 - benchmarks.report.summary - INFO:
===== Foretoken Benchmark Result =====
  Model      : Qwen3.6-27B
  Requests   : 20
  Parallel    : 4
  Rate       : 5.0 req/s
  Success     : 20/20 (100.00%)
  Latency     mean=1.6679s p50=1.6494s p95=1.7332s p99=1.9512s
  TTFT        mean=0.1937s p50=0.1872s p95=0.2205s p99=0.3949s
  TPOT        mean=0.0234s p50=0.0235s p95=0.0249s p99=0.0249s
  Request/s   : 2.1755
  Token/s     : 139.2326
  Tok/s/user  : 34.8081
  Wall time   : 9.1933s
=====
2026-08-10 16:03:27,235 - benchmarks.runner.base - INFO: Results saved: results/20260810_160
309

```


Charts 11 1-6 of 11




Output Throughput (tok 1)




Req Throughput (req 1)




Request Rate (req 1)



Total Throughput (tok 1)





HF数据集






```
foretoken bench \  
  --url http://127.0.0.1:8008/v1/chat/completions \  
  --model Qwen3.6-27B \  
  --dataset weijiezz/foretoken-trace:conversation \  
  --parallel 4 \  
  --number 20 \  
  --wandb
```


Q Search panels with regex

🖨️ ⋮ ⚙️ Settings + New report + Add panels


⋮ ▾ Charts 11

1-6 of 11 < > 🔍 ⚙️ ⋮




⋮ ▾ Output Throughput (tok) 1

🔍 ⚙️ ⋮




⋮ ▾ Req Throughput (req) 1

🔍 ⚙️ ⋮




⋮ ▾ Request Rate (req) 1


🔍 ⚙️ ⋮



⋮ ▾ Total Throughput (tok) 1

🔍 ⚙️ ⋮





本地数据集

```
foretoken bench \
--url http://127.0.0.1:8008/v1/chat/completions \
--model Qwen3.6-27B \
--dataset /home/wshiah/code/zhuting/foretoken/conversation.jsonl \
--parallel 4 \
--number 20 \
--wandb
```

```

• (eval) wshiah@zxcpu4:~/code/zhuting/bench/foretoken/benchmarks$ foretoken bench \
  --url http://127.0.0.1:8008/v1/chat/completions \
  --model Qwen3.6-27B \
  --dataset /home/wshiah/code/zhuting/foretoken/conversation.jsonl \
  --parallel 4 \
  --number 20 \
  --wandb
2026-08-10 16:12:13,299 INFO [benchmarks.main]
=====
Foretoken Benchmark
=====
Configuration:
URL      : http://127.0.0.1:8008/v1/chat/completions
Model    : Qwen3.6-27B
Parallel : [4]
Number   : [20]
Rate     : INF (no pacing)
Open Loop : False
Stream   : True
Dataset  : /home/wshiah/code/zhuting/foretoken/conversation.jsonl

2026-08-10 16:12:20,042 INFO [benchmarks.report.wandb_logger] W&B logging enabled: project=f
oretoken-bench name=Qwen3.6-27B_20260810_161213 group=-
Benchmarking: 100%|████████████████████████████████████████| 20/20 [01:09<00:00, 3.46s/it]
2026-08-10 16:13:29,348 INFO [benchmarks.runner.base] Benchmark finished!
2026-08-10 16:13:29,353 INFO [benchmarks.report.summary]
===== Foretoken Benchmark Result =====
Model      : Qwen3.6-27B
Requests   : 20
Parallel    : 4
Success     : 20/20 (100.00%)
Latency     mean=13.7707s p50=12.4209s p95=21.8120s p99=24.0749s
TTFT        mean=5.8892s p50=4.6751s p95=12.6970s p99=14.8740s
TPOT        mean=0.0621s p50=0.0626s p95=0.0810s p99=0.0884s
Request/s   : 0.2886
Token/s     : 36.9421
Tok/s/user  : 9.2355
Wall time   : 69.2976s
=====





```


Search panels with regex


Settings + New report + Add panels

Charts 11


1-6 of 11




Output Throughput (tok 1)




Req Throughput (req 1)




Request Rate (req 1)



Total Throughput (tok 1)





多数据集

```
foretoken bench \
--url http://127.0.0.1:8008/v1/chat/completions \
--model Qwen3.6-27B \
--dataset r0b0t1ab/qwen3.8-max-distillation-50k:train,ianncity/GLM-5.2-Conversation:train \
--parallel 4 \
--number 20 \
--wandb
```

```

• (eval) wshiah@zxcpu4:~/code/zhuting/bench/foretoken/benchmarks$ foretoken bench \
  --url http://127.0.0.1:8008/v1/chat/completions \
  --model Qwen3.6-27B \
  --dataset r0b0t1ab/qwen3.8-max-distillation-50k:train,ianncity/GLM-5.2-Conversation:train \
  --parallel 4 \
  --number 20 \
  --wandb
2026-08-10 16:14:32,902 INFO [benchmarks.main]
=====
Foretoken Benchmark
=====
Configuration:
  URL       : http://127.0.0.1:8008/v1/chat/completions
  Model      : Qwen3.6-27B
  Parallel   : [4]
  Number     : [20]
  Rate       : INF (no pacing)
  Open Loop  : False
  Stream     : True
  Dataset    : ['r0b0t1ab/qwen3.8-max-distillation-50k:train', 'ianncity/GLM-5.2-Conversation:train'] (total number across all)

2026-08-10 16:14:32,903 INFO [benchmarks.runner.multi_dataset] Dataset 1/2: r0b0t1ab/qwen3.8-max-distillation-50k:train (number=10)
2026-08-10 16:15:17,878 INFO [benchmarks.report.wandb_logger] W&B logging enabled: project=foretoken-bench name=Qwen3.6-27B_20260810_161432_00_r0b0t1ab_qwen3.8-max-distillation-50k_train group=Qwen3.6-27B_20260810_161432
Benchmarking: 100%|████████████████████████████████████████| 10/10 [00:09<00:00, 1.03it/s]
2026-08-10 16:15:27,557 INFO [benchmarks.runner.base] Benchmark finished!
2026-08-10 16:15:27,561 INFO [benchmarks.report.summary]
===== Foretoken Benchmark Result =====
  Model      : Qwen3.6-27B
  Requests   : 10
  Parallel    : 4
  Dataset     : r0b0t1ab/qwen3.8-max-distillation-50k:train
  Success     : 10/10 (100.00%)
  Latency     : mean=3.1768s p50=3.1509s p95=3.4006s p99=3.4969s
  TTFT        : mean=0.3476s p50=0.3675s p95=0.6074s p99=0.6876s
  TPOT        : mean=0.0223s p50=0.0219s p95=0.0240s p99=0.0242s
  Request/s   : 1.0342
  Token/s     : 132.3736
  Tok/s/user  : 33.0934
  Wall time   : 9.6696s
=====
2026-08-10 16:15:29,289 INFO [benchmarks.runner.multi_dataset] Dataset 2/2: ianncity/GLM-5.2-Conversation:train (number=10)
2026-08-10 16:15:46,271 INFO [benchmarks.report.wandb_logger] W&B logging enabled: project=foretoken-bench name=Qwen3.6-27B_20260810_161432_01_ianncity_GLM-5.2-Conversation_train group=Qwen3.6-27B_20260810_161432
Benchmarking: 100%|████████████████████████████████████████| 10/10 [00:10<00:00, 1.00s/it]
2026-08-10 16:15:56,316 INFO [benchmarks.runner.base] Benchmark finished!
2026-08-10 16:15:56,321 INFO [benchmarks.report.summary]
===== Foretoken Benchmark Result =====
  Model      : Qwen3.6-27B
  Requests   : 10
  Parallel    : 4
  Dataset     : ianncity/GLM-5.2-Conversation:train
  Success     : 10/10 (100.00%)
  Latency     : mean=3.3797s p50=3.4189s p95=3.4785s p99=3.4800s

```

TTFT mean=0.4464s p50=0.4398s p95=0.6884s p99=0.6899s
TPOT mean=0.0231s p50=0.0225s p95=0.0255s p99=0.0260s
Request/s : 0.9964
Token/s : 127.5385
Tok/s/user : 31.8846
Wall time : 10.0362s

=====

2026-08-10 16:15:58,320 INFO [benchmarks.report.summary]

===== Foretoken Benchmark Result =====

Model : Qwen3.6-27B
Requests : 20
Parallel : 4
Datasets : ['r0b0tlab/qwen3.8-max-distillation-50k:train', 'ianncity/GLM-5.2-Conversation:train']
Success : 20/20 (100.00%)
Latency mean=3.2783s p50=3.2436s p95=3.4824s p99=3.5132s
TTFT mean=0.3970s p50=0.3773s p95=0.6911s p99=0.7043s
TPOT mean=0.0227s p50=0.0219s p95=0.0249s p99=0.0258s
Request/s : 1.0149
Token/s : 129.9110
Tok/s/user : 32.4778
Wall time : 19.7058s

=====

2026-08-10 16:15:58,326 INFO [benchmarks.runner.base] Results saved: results/20260810_161432

🔍 Search runs



🔍 1 filter









🕒 NAME 2 listed

🔴 Qwen3.6-27B_...versation_train


🟢 Qwen3.6-27B_...ation-50k_train

🔍 Charts 11


1-6 of 11 < > 🔗 ⚙️ ⋮




🔍 Output Throughput (tok 1




🔍 Req Throughput (req 1



🔍 Request Rate (req 1



🔍 Total Throughput (tok 1



🔍 System 22

1-6 of 22 < > 🔗 ⚙️ ⋮

